

Lecture 10: Linear regression with R

- Linear regression
 - Simple linear regression
 - Multiple linear regression
- Analysis of variance (ANOVA)

10.1 Linear regression

10.1.1 Simple linear regression

I will use simulated data for this section. This will help you understand model assumptions.

For this lecture, try to understand the data generating process.

Some background information.

- I spent a few weeks shopping for a diamond ring.
- I browsed many diamonds with G or H color, VVS1 or VVS2 clarity, excellent cut and between 1.00 and 1.50 carats.
 - I found out that the prices of such diamonds were predictable.
 - My gut feeling was that there was an algorithm for pricing, which waw then rounded it to the nearest 100x.
- As a statistician, I was thinking of fitting a linear regression to explore the relationship between the 4C's and the price.
 - So that I could pick the one that was significantly cheaper than it should be according to my statistical model.
 - Therefore having the best price-quality ratio.
 - I was too young and naive.
- **The most important thing is carat!**
- But still the other 3C's and some other minor factors may affect the price by a smaller margin.
 - G color is more expensive, H color is less expensive.

The assumptions on the small effects are,

- Normal distribution.
 - This small deviation is **normally distributed with mean 0**.
- **Constant variance**.
 - This small deviation is the same for all diamonds, regardless of carat.

Therefore the data follows,

$$price_i = \beta_0 + \beta_1 \times carat_i + \epsilon_i$$

where $\epsilon_i \sim_{i.i.d} Normal(0, \sigma^2)$ is the small deviation, which is sometimes called the noise.

In [1]:

```
# Set seed for reproducibility
set.seed(613)

# 100 diamond rings ranging from 1 to 1.5 carats
n <- 100
carat <- runif(n = n, min = 1, max = 1.5)
color <- as.factor(sample(c("G", "H"), size = n, replace = TRUE)
)


# Normally distributed deviation from my perfect model caused by
other factors
error <- rnorm(n = n, mean = 0, sd = 0.5)

# On average:
# 1.00 carat: $10k;
# .....;
# 1.50 carat: $20k;

# Comparing to the average price.
# G color increases the price by 500
# H color decreases the price by 500
price <- -5 + 20 * carat + error + ifelse(color == "G", 0.5, -0.
5)

plot(x = carat, y = price, main = "Price vs Carat - Fake Data")
```

Price vs Carat - Fake Data



The model is,

$$\widehat{price}_i = \hat{\beta}_0 + \hat{\beta}_1 \times carat_i$$

The goal is to de-noise and capture the association - find a straight line that best describes the association. How?

Consider the residual (e_i), which is the deviation of the observed data from the best fit line.

$$e_i = price_i - \widehat{price}_i = price_i - (\hat{\beta}_0 + \hat{\beta}_1 \times carat_i)$$

It makes sense to find a line that minimizes the sum of the residuals squared $\sum_{i=1}^n e_i^2$.

Why squared?

The math is,

in EPIB 621.

There is closed-form solution to linear model coefficients - nice.

Fit a linear model of the data considering only carat

In [2]:

```
fit1 <- lm(price~carat); fit1
```

Call:

```
lm(formula = price ~ carat)
```

Coefficients:

(Intercept)	carat
-6.159	20.903

In [3]:

```
# If our data is a data.frame, note the data argument.
diamond <- data.frame(Price = price, Carat = carat, Color = color)
head(diamond)
aggregate(Price~Color, data = diamond, mean)
table(diamond$Color)
```

A data.frame: 6 × 3

Price	Carat	Color
<dbl>	<dbl>	<fct>
24.95654	1.489319	G
23.25550	1.477755	H
15.57460	1.051832	H
21.77606	1.411278	H
18.21557	1.162526	G
24.53011	1.465442	G

A data.frame: 2 × 2

Color	Price
<fct>	<dbl>
G	21.54011
H	19.09459
G	H
48	52

In [4]:

```
fit1 <- lm(Price~Carat, data = diamond)
summary(fit1)
```

Call:

```
lm(formula = Price ~ Carat, data = diamond)
```

Residuals:

Min	1Q	Median	3Q	Max
-1.56430	-0.56169	-0.02658	0.60028	1.82743

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-6.1591	0.6008	-10.25	<2e-16 ***
Carat	20.9027	0.4715	44.34	<2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.'
0.1 ' ' 1


Residual standard error: 0.7509 on 98 degrees of freedom

Multiple R-squared: 0.9525, Adjusted R-squared:
0.952

F-statistic: 1966 on 1 and 98 DF, p-value: < 2.2e-
16

In [5]:

```
plot(x = carat, y = price, main = "Price vs Carat - Fake Data")  
# Draw the best fit line from our linear model.  
abline(fit1, col = "red")
```



In [6]:

```
# Explore the structure of fit1 and summary(fit1) to see how to  
extract values.  
# str(fit1)  
# str(summary(fit1))
```

Interpretation of the coefficients along with confidence intervals of the estimated coefficients.

In [7]:

```
summary(fit1)$coefficients
```

A matrix: 2 × 4 of type dbl

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-6.159122	0.6007728	-10.25200	3.437439e-17
Carat	20.902675	0.4714502	44.33697	1.167533e-66

In [8]:

```
# Confidence intervals - default is 95%  
confint(fit1, level = 0.95)
```

A matrix: 2 × 2 of type dbl

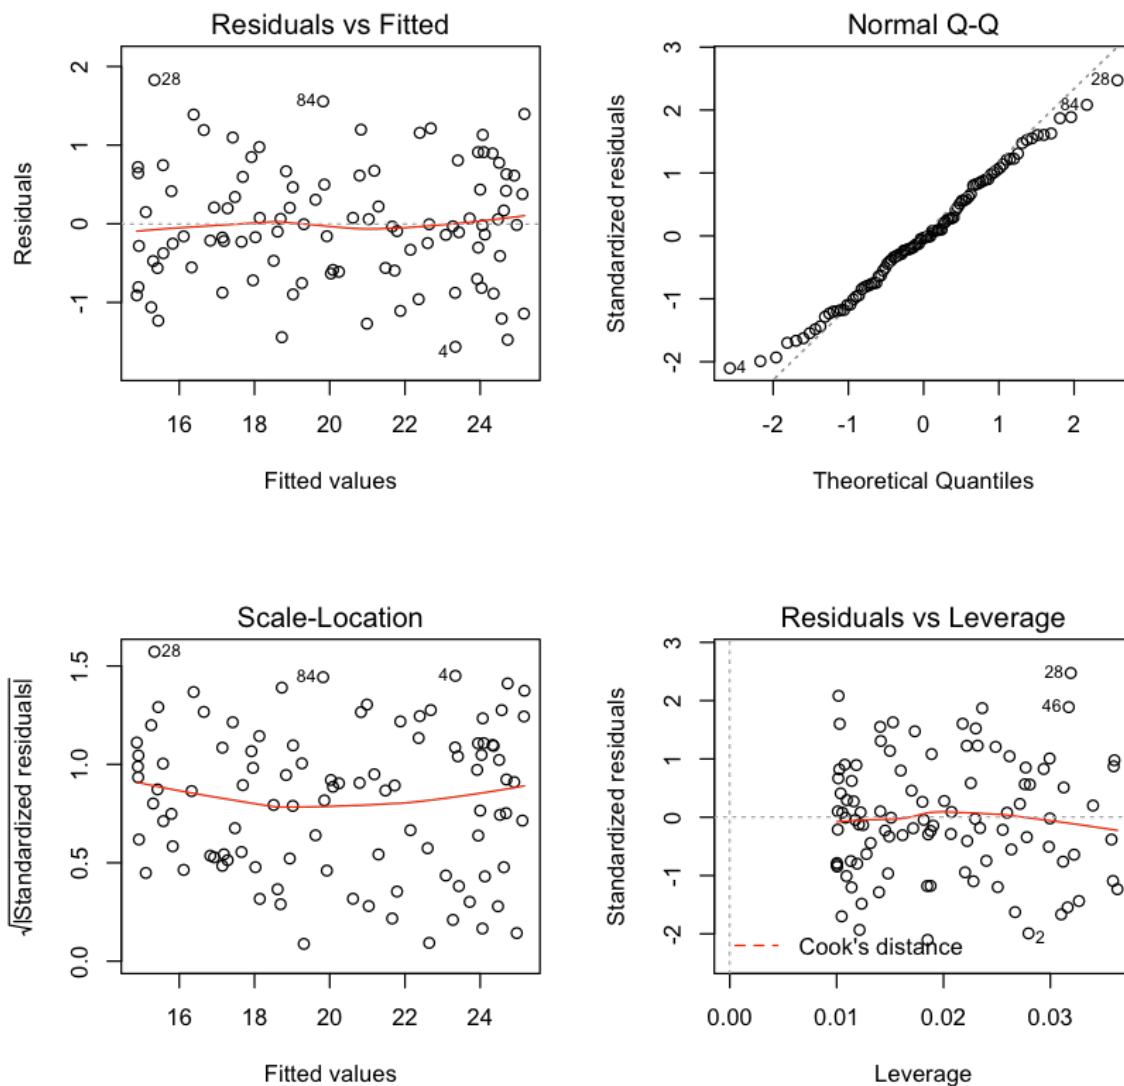
	2.5 %	97.5 %
(Intercept)	-7.351336	-4.966908
Carat	19.967097	21.838252

Check model fit and assumptions using plots

- First plot checks the constant variance assumption
- Second plot checks the normality assumption

In [9]:

```
par(mfrow = c(2,2))  
plot(fit1)
```



10.1.2 Multiple linear regression

Simply more than one x 's.

Fit a linear model considering both carat and color

In [10]:

```
fit2 <- lm(Price~Carat+Color, data = diamond)
summary(fit2)
```

Call:

```
lm(formula = Price ~ Carat + Color, data = diamond)
```

Residuals:

Min	1Q	Median	3Q	Max
-0.94938	-0.42954	-0.02353	0.34489	1.09171

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-4.6503	0.4227	-11.00	<2e-16 ***
Carat	20.1772	0.3209	62.88	<2e-16 ***
ColorH	-1.1376	0.1023	-11.12	<2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.'
0.1 ' ' 1

Residual standard error: 0.5004 on 97 degrees of freedom

Multiple R-squared: 0.9791, Adjusted R-squared:
0.9787

F-statistic: 2275 on 2 and 97 DF, p-value: < 2.2e-16

In [11]:

```
# Adjust for all other variables in the dataset other than Price
.  
# lm(Price~., data = diamond)  
  
# Adjust for interaction terms - effect measure modification EPI  
B 603.  
# lm(Price~ Carat + Color + Carat:Color, data = diamond)  
# Equivalent to  
# lm(Price~ Carat * Color, data = diamond)
```

Prediction using the regression model

Predict the average price of an 1.22 carat diamond with G color - confidence interval.

In [12]:

```
predict(fit2, newdata = data.frame(Carat=1.22, Color="G"), interval = "confidence")
```

A matrix: 1 × 3 of type dbl

fit	lwr	upr
19.96585	19.81413	20.11758

Predict the price of an 1.22 carat diamond with G color - prediction interval.

In [13]:

```
predict(fit2, newdata = data.frame(Carat=1.22, Color="G"), interval = "prediction")
```

A matrix: 1 × 3 of type dbl

fit	lwr	upr
19.96585	18.9611	20.97061

Draw confidence interval and prediction interval for simple linear regression (fit1)

Hard to visualize when there are more than one x 's.

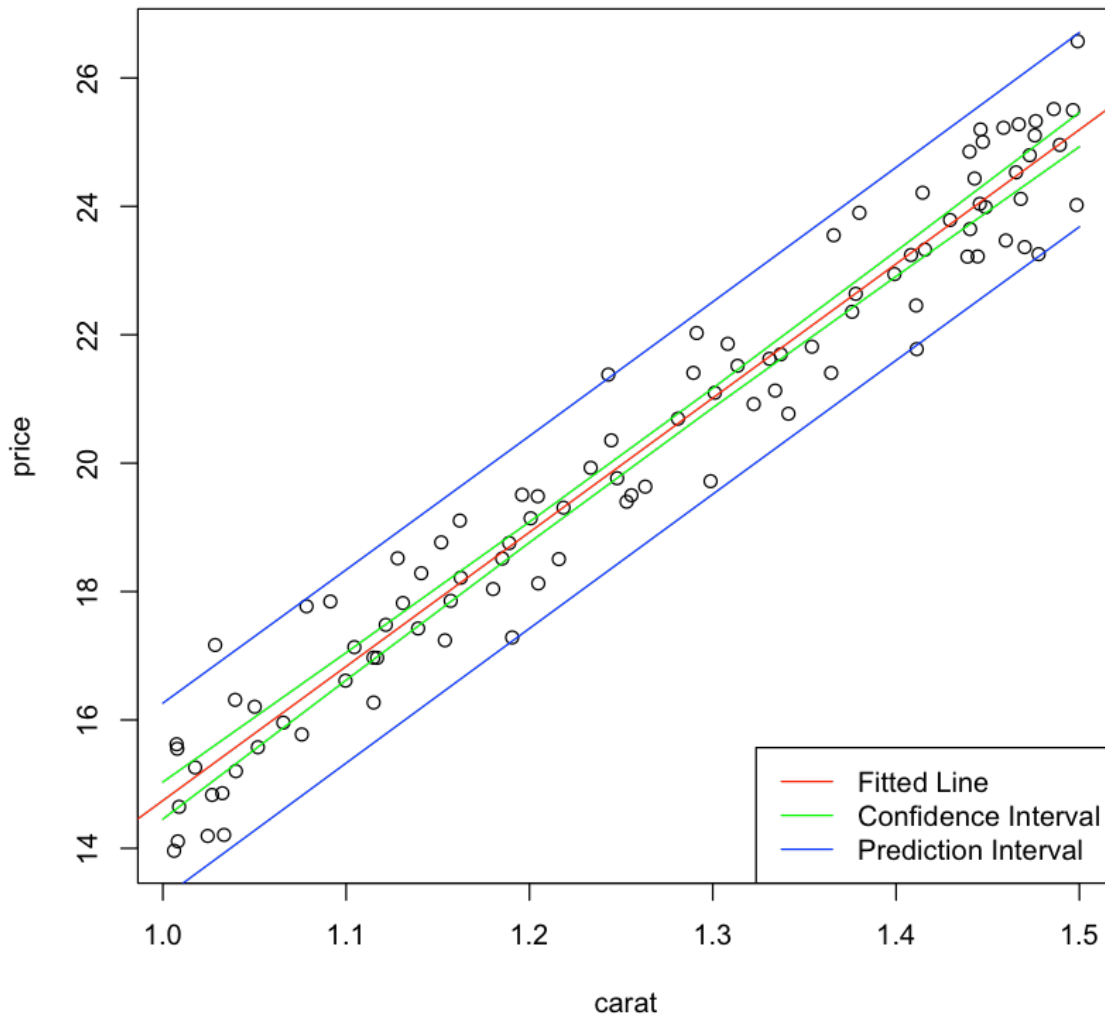
In [14]:

```
CI <- predict(fit1, newdata = data.frame(Carat=seq(1,1.5,0.1)),
interval = "confidence")
PI <- predict(fit1, newdata = data.frame(Carat=seq(1,1.5,0.1)),
interval = "prediction")
plot(x = carat, y = price, main = "Price vs Carat - Fake Data")
# Draw the best fit line from our linear model.
abline(fit1, col = "red")
lines(x = seq(1,1.5,0.1), y = CI[,2], col = "green")
lines(x = seq(1,1.5,0.1), y = CI[,3], col = "green")

lines(x = seq(1,1.5,0.1), y = PI[,2], col = "blue")
lines(x = seq(1,1.5,0.1), y = PI[,3], col = "blue")

legend("bottomright", col = c("red", "green", "blue"), lty = c(1
,1,1),
      legend = c("Fitted Line", "Confidence Interval", "Predict
ion Interval"))
```

Price vs Carat - Fake Data



Programming-wise, simple linear regression and multiple linear regressions are the same.

10.2 Analysis of Variance

10.2.1 ANOVA without regression

Recall that we could use t-test to test whether two treatment effects are significantly different.

In [15]:

```
price.G <- subset(x = diamond, Color == "G")$Price
price.H <- subset(x = diamond, Color == "H")$Price
mu.G <- mean(price.G)
mu.H <- mean(price.H)
mu.tot <- mean(diamond$Price)
data.frame(G=mu.G, H=mu.H, Total=mu.tot)
```

A data.frame: 1 × 3

G	H	Total
<dbl>	<dbl>	<dbl>
21.54011	19.09459	20.26844

In [16]:

```
t.test(price.G, price.H)
```

Welch Two Sample t-test

```
data: price.G and price.H
t = 3.8049, df = 97.895, p-value = 0.0002472
alternative hypothesis: true difference in means is
not equal to 0
95 percent confidence interval:
 1.170011 3.721021
sample estimates:
mean of x mean of y
 21.54011  19.09459
```

This is a direct comparison of the two means.

- G color is significantly more expensive than H color.

We can also use the variance to test whether two groups are different. How?

- Decompose the total variation in the dataset into within group variation and between group variation.
- Think of the test as: if too much of the total variation comes from the between group variation, we believe that the two groups are different.

In [17]:

```
myanova <- aov(Price~Color, data = diamond)
summary(myanova)
```

```
          Df Sum Sq Mean Sq F value    Pr(>F)
Color         1   149.3   149.27    14.42 0.000254 ***
Residuals    98 1014.4    10.35
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.'
0.1 ' ' 1
```

The null hypothesis that the two groups are equal is rejected.

In [18]:

```
# Reproduce the test result from scratch
# ss means sum of squares
ss.within = sum((price.G-mu.G)^2) + sum((price.H-mu.H)^2)
ss.between = (mu.G-mu.tot)^2 * 48 + (mu.H-mu.tot)^2 * 52
ss.total = sum((price-mu.tot)^2)
data.frame(Within = ss.within, Between = ss.between, Total = ss.
total)
```

A data.frame: 1 × 3

Within	Between	Total
<dbl>	<dbl>	<dbl>
1014.435	149.2745	1163.709

10.2.2 ANOVA in regression

- We saw that the calculations involve within group variance calculation for each group.
- We had 2 groups. What if we had more?
- What if we had many more? e.g. carat in the first dataset - a continuous variable.

In [19]:

```
head(diamond$Carat, 20)
```

1.48931883147452	1.4777549788123	1.05183249094989
1.41127772070467	1.16252580692526	1.46544234571047
1.28101023298223	1.44567985611502	1.09136897884309
1.21604923321865	1.30815492232796	1.18015707843006
1.20065634301864	1.11689501185901	1.29879072017502
1.35407601017505	1.11498201265931	1.13925164530519
1.01760521158576	1.03983685944695	

The logic: fit a linear model, decompose the total variation into model variation (between group) and residual variation (within group).

In [20]:

```
anova(fit1)
```

A anova: 2 × 5

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
	<int>	<dbl>	<dbl>	<dbl>	<dbl>
Carat	1	1108.44923	1108.4492293	1965.767	1.167533e-66
Residuals	98	55.25986	0.5638761	NA	NA

In [21]:

```
anova(fit2)
```

A anova: 3 × 5

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
	<int>	<dbl>	<dbl>	<dbl>	<dbl>
Carat	1	1108.44923	1108.4492293	4426.0268	9.597565e-83
Color	1	30.96729	30.9672870	123.6521	5.228797e-19
Residuals	97	24.29257	0.2504389	NA	NA

Compare two regression models.

- Our fit1 is a reduced version of fit2
- Test whether the deletion of "Color" significantly affect model fit.
 - Rejected: use full model.

In [22]:

```
print(anova(fit1, fit2))
```

Analysis of Variance Table

Model 1: Price ~ Carat

Model 2: Price ~ Carat + Color

	Res.Df	RSS	Df	Sum of Sq	F	Pr(>F)
1	98	55.260				
2	97	24.293	1	30.967	123.65	< 2.2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.'
0.1 ' ' 1

Assignment 2